

# COMMUNICATING RESULTS

*Sri Kanajan*

---

## COMMUNICATING RESULTS

---

# LEARNING OBJECTIVES

- Explain the trade-offs between the precision and recall of a model while articulating the cost of false positives vs. false negatives
- Describe the difference between visualization for presentations vs. exploratory data analysis
- Identify the components of a concise, convincing report and how they relate to specific audiences/stakeholders

---

**COURSE**

---

**PRE-WORK**

---

## **PRE-WORK REVIEW**

---

- Understand results from a confusion matrix and measure true positive rate and false positive rate
- Create and interpret results from a binary classification problem

---

**OPENING**

---

# COMMUNICATING RESULT

---

## **WE BUILT A MODEL! NOW WHAT?**

---

- We've built our model, but there is still a gap between your Notebook with plots/figures and a slideshow needed to present your results.
- Classes so far have focused on two core concepts:
  - developing consistent practices
  - interpreting metrics to evaluate and improve model performance
- But what does that mean to your audience?

---

## WE BUILT A MODEL! NOW WHAT?

---

- Imagine how a non-technical audience might respond to the following statements:
  - The predictive model I built has an accuracy of 80%.
  - Logistic regression was optimized with L2 regularization.
  - Gender was more important than age in the predictive model because it has a larger coefficient.
  - Here's the AUC chart that shows how well the model did.

---

## WE BUILT A MODEL! NOW WHAT?

---

- Who is your audience? Are they technical? What are their concerns?
- Remember: in a business setting, you may be *the only person* who can interpret what you've built.
- Some people may be familiar with basic visualization, but you will likely have to do a lot of “hand holding”.
- You need to be able to efficiently explain your results in a way that makes sense to **all** stakeholders (technical or not).



---

## **WE BUILT A MODEL! NOW WHAT?**

---

- Today, we'll focus on communicating results for “simpler” problems, but this applies to any type of model you may work with.
- First, let's review classification metrics, review our knowledge, and talk about how we might communicate what we know.

---

**REVIEW**

---

# BACK TO THE CONFUSION MATRIX

---

## BACK TO THE CONFUSION MATRIX

---

- Confusion matrices allow for the interpretation of correct and incorrect predictions for *each class label*.
- It is the first step for the majority of classification metrics and goes deeper than just accuracy.

		<u>True class</u>	
		<b>p</b>	<b>n</b>
<u>Hypothesized class</u>	<b>Y</b>	True Positives	False Positives
	<b>N</b>	False Negatives	True Negatives

$$\text{fp rate} = \frac{FP}{N}$$
$$\text{tp rate} = \frac{TP}{P}$$
$$\text{precision} = \frac{TP}{TP+FP} \quad \text{recall} = \frac{TP}{P}$$
$$\text{accuracy} = \frac{TP+TN}{P+N}$$
$$\text{F-measure} = \frac{2}{1/\text{precision} + 1/\text{recall}}$$

**Column totals:**

# THE MATH FOR RECALL

- Recall is the count of predicted *true positives* over the total count of that class label.
- This is the same as True Positive Rate or *sensitivity*.

		<u>True class</u>			
		<b>p</b>	<b>n</b>		
<u>Hypothesized class</u>	<b>Y</b>	True Positives	False Positives	$fp\ rate = \frac{FP}{N}$	$tp\ rate = \frac{TP}{P}$
	<b>N</b>	False Negatives	True Negatives	$precision = \frac{TP}{TP+FP}$	$recall = \frac{TP}{P}$
Column totals:		<b>P</b>	<b>N</b>	$accuracy = \frac{TP+TN}{P+N}$	$F\text{-measure} = \frac{2}{1/precision + 1/recall}$

# THE MATH FOR PRECISION

- Precision, or positive predicted value, is calculated as the count of predicted true positives over the count of all values predicted to be positive.

		<u>True class</u>			
		<b>p</b>	<b>n</b>		
<u>Hypothesized class</u>	<b>Y</b>	True Positives	False Positives	$fp\ rate = \frac{FP}{N}$	$tp\ rate = \frac{TP}{P}$
	<b>N</b>	False Negatives	True Negatives	$precision = \frac{TP}{TP+FP}$	$recall = \frac{TP}{P}$
Column totals:		<b>P</b>	<b>N</b>	$accuracy = \frac{TP+TN}{P+N}$	
				$F\text{-measure} = \frac{2}{1/precision + 1/recall}$	

**DEMO**

---

# UNDERSTANDING TRADEOFF

---

# UNDERSTANDING TRADEOFF

---

- Let's consider the following data problem: we are given a data set in order to predict or identify traits for typically late flights.
- Optimizing toward recall, we could assume that every flight will be delayed.
- The trade-off, a lower precision, is that this could create even further delays, missed flights, etc.



---


## UNDERSTANDING TRADEOFF

---

- Optimizing toward precision, we would specifically look to identify flights that will be late.
- The trade-off here would be lower recall. We might miss flights that would be delayed, causing a strain on the system.

# PRECISION RECALL CURVE

- ▶ Another way to visualize a model and compare between models



---

## INTRODUCTION

---

# SHOWING WORK

---

## SHOWING WORK

---

- We've spent a lot of time exploring our data and building a reasonable model that performs well.
- However, if we look at our visuals, they are most likely:
  - Statistically heavy: Most people don't understand histograms.
  - Overly complicated: Scatter matrices produce too much information.
  - Poorly labeled: Code doesn't require adding labels, so you may not have added them.

---

# SHOWING WORK

---

- In order to convey important information to our audience, make sure our charts are:
  - Simplified
  - Easily interpretable
  - Clearly labeled

---

## **SIMPLIFIED**

---

- At most, you'll want to include figures that either explain a variable on its own or explain that variable's relationship with a target.
- If your model used a data transformation (like natural log), just visualize the original data.
- Try to remove any unnecessary complexity.

---

## **EASILY INTERPRETABLE**

---

- Any stakeholder looking at a figure should be seeing the exact same thing you're seeing.
- A good test for this is to share the visual with others less familiar with the data and see if they come to the same conclusion.
- How long did it take them?

---

## **CLEARLY LABELED**

---

- Take the time to clearly label your axis, title your plot, and double check your scales - especially if the figures should be comparable.
- If you're showing two graphs side by side, they should follow the same Y axis.



---

## QUESTION TO ASK

---

- When building visuals for another audience, ask yourself these questions:
  - **Who:** Who is my target audience for the visual?
  - **What:** What do they already know about this project? What do they need to know?
  - **How:** How does my project affect this audience? How might they interpret (or misinterpret) the data?

**DEMO**

---

# VISUALIZING MODELS OVER VARIABLES

---

## **VISUALIZING MODELS OVER VARIABLES**

---

- One effective way to explain your model over particular variables is to plot the predicted values against the most explanatory variables.
- For example, in logistic regression, plotting the probability of a class against a variable can help explain the range of effect of the model.

---

# **VISUALIZING MODELS OVER VARIABLES**

---

- We'll use the flight delay data for all following examples. Let's build our first model and plot.
- Open the starter code from the class repo and follow along.

---

# VISUALIZING MODELS OVER VARIABLES

---

```
# read in the file and generate a quick model (assume we've done the data exploration already)
```

```
import pandas as pd
import sklearn.linear_model as lm
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('../assets/dataset/flight_delays.csv')
```

```
df = df.join(pd.get_dummies(df['DAY_OF_WEEK'], prefix='dow'))
```

```
df = df[df.DEP_DEL15.notnull()].copy()
```

---

# VISUALIZING MODELS OVER VARIABLES

---

```
# Build a model
model = lm.LogisticRegression()
features = ['dow_1', 'dow_2', 'dow_3', 'dow_4', 'dow_5', 'dow_6']
model.fit(df[features + ['CRS_DEP_TIME']], df['DEP_DEL15'])

df['probability'] = model.predict_proba(df[features + ['CRS_DEP_TIME']]).T[1]
```

---

# VISUALIZING MODELS OVER VARIABLES


---

```
# Create a plot
ax = plt.subplot(111)
colors = ['blue', 'green', 'red', 'purple', 'orange', 'brown']
for e, c in enumerate(colors):
    df[df[features[e]] == 1].plot(x='CRS_DEP_TIME', y='probability',
kind='scatter', color = c, ax=ax)

ax.set(title='Probability of Delay\n Based on Day of Week and Time of Day')
```

# VISUALIZING MODELS OVER VARIABLES

- This visual can help showcase the range of effect on delays from both day of the week and time of day.
- Given this model, some days are more likely to have delays than others.
- The likelihood of delay increases as the day goes on.





---

# ACTIVITY: TRY IT OUT

---



## EXERCISE

### DIRECTIONS

1. Adjust the model to make delay predictions using airlines instead of day of week, and time, then plot the effect on `CRS_DEP_TIME=1`.
2. Try plotting the inverse: pick either model and plot the effect on `CRS_DEP_TIME=0`.

### DELIVERABLE

The new plots

**DEMO**

---

# **VISUALIZING PERFORMANCE AGAINST BASELINE**

---

## **VISUALIZING PERFORMANCE AGAINST BASELINE**


---

- Another approach of visualization is the effect of your model against a baseline, or - even better - against previous models.
- Plots like this will also be useful when talking to your peers - other data scientists or analysts who are familiar with your project and interested in the progress you've made.

# VISUALIZING PERFORMANCE AGAINST BASELINE


▸ For classification, we've practiced plotting AUC and precision-recall plots. Consider the premise of each:

- AUC plots explain and represent “accuracy” as having the largest area under the curve. Good models will be high and to the left.
- For precision-recall plots, it will depend on the *cost* requirements. Either a model will have good recall at the cost of precision or vice versa.



# VISUALIZING PERFORMANCE AGAINST BASELINE

- When comparing multiple models:
  - For AUC plots, you'll be interested in which model has the *largest* area under the curve.
  - For precision-recall plots, based on the cost requirement, you are looking at which model has the best precision given the same recall, or the best recall given the same precision.



---

# VISUALIZING PERFORMANCE AGAINST BASELINE

---

- Follow along with the starter code located in the class repo.
- We've plotted several models for AUC: a dummy model and additional features.

```
model0 = dummy.DummyClassifier()  
model0.fit(df[features[1:-1]], df.DEP_DEL15)  
df['probability_0'] = model0.predict_proba(df[features[1:-1]]).T[1]
```

```
model = lm.LogisticRegression()  
model.fit(df[features[1:-1]], df.DEP_DEL15)  
df['probability_1'] = model.predict_proba(df[features[1:-1]]).T[1]
```

---

# VISUALIZING PERFORMANCE AGAINST BASELINE

---


```
ax = plt.subplot(111)
vals = metrics.roc_curve(df.DEP_DEL15, df.probability_0)
ax.plot(vals[0], vals[1])
vals = metrics.roc_curve(df.DEP_DEL15, df.probability_1)
ax.plot(vals[0], vals[1])
```

```
ax.set(title='Area Under the Curve for prediction delayed=1', ylabel='TRP',
xlabel='FRP', xlim=(0, 1), ylim=(0, 1))
```

# VISUALIZING PERFORMANCE AGAINST BASELINE

▸ This plot showcases:

1. The model using data outperforms a baseline dummy model.
2. By adding other features, there's some give and take with probability as the model gets more complicated.





---

# ACTIVITY: TRY IT OUT

---



## EXERCISE

### DIRECTIONS

1. In a similar approach, use the sklearn `precision_recall_curve` function to enable you to plot the precision-recall curve of the four models from above. Keep in mind precision in the first array is returned from the function, but the plot shows it as the y-axis.
2. Explain what is occurring when the recall is below 0.2.
3. Based on this performance, is there a clear winner at different thresholds?
4. **Bonus:** Redo both the AUC and precision-recall curves using models that have been cross validated using `kfold`. How do these new figures change your expectations for performance?

### DELIVERABLE

The new plots and associated answers

---

**INDEPENDENT PRACTICE**

---

**PROJECT PRACTICE**

---

# ACTIVITY: PROJECT PRACTICE

---



## EXERCISE

### DIRECTIONS (45 minutes)

Using models built from the flight data problem earlier in class, work through the same problems. Your data and models should already be accessible. Your goals:

1. There are *many* ways to manipulate this data set. Perhaps mix both carrier and day of week in your model. Consider what is a proper "categorical" variable, and keep *only* what is significant. You will easily have 20+ variables. Aim to have a visual that clearly explain the relationship of variables you've used against the predicted flight delay.
2. Generate the AUC or precision-recall curve (based on which you think makes more sense), and have a statement that defines, compared to a baseline, how your model performs and any caveats. For example: "My model on average performs at x rate, but the features under-perform and explain less of the data at these thresholds." Consider this as practice for your own project, since the steps you'll take to present your work will be relatively similar.

### DELIVERABLE

New models and performance statement

---

**CONCLUSION**

---

# TOPIC REVIEW

---

## REVIEW AND NEXT STEPS

---

- What do precision and recall mean? How are they similar and different to True Positive Rate and False Positive Rate?
- How does cost benefit analysis play a role in building models?
- What are at least two very important details to consider when creating visuals for a project's stakeholders?
- Why would an AUC plot work well for a data science audience but not for a business audience? What would be a more effective visualization for that group?

**COURSE**

---

**BEFORE NEXT CLASS**

## **LESSON**

---

# **EXIT TICKET**

**DON'T FORGET TO FILL OUT YOUR EXIT TICKET**